

logistic-regression-assignment

September 26, 2024

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[2]: data = pd.read_csv(r"C:\\Users\\prana\\Documents\\Python data\\SeattleWeather.  
csv")  
data
```

```
[2]:
```

	DATE	PRCP	TMAX	TMIN	RAIN
0	01-01-1948	0.47	51	42	True
1	02-01-1948	0.59	45	36	True
2	03-01-1948	0.42	45	35	True
3	04-01-1948	0.31	45	34	True
4	05-01-1948	0.17	45	32	True
...
25546	10-12-2017	0.00	49	34	False
25547	11-12-2017	0.00	49	29	False
25548	12-12-2017	0.00	46	32	False
25549	13-12-2017	0.00	48	34	False
25550	14-12-2017	0.00	50	36	False

[25551 rows x 5 columns]

```
[3]: data.head(5)
```

```
[3]:
```

	DATE	PRCP	TMAX	TMIN	RAIN
0	01-01-1948	0.47	51	42	True
1	02-01-1948	0.59	45	36	True
2	03-01-1948	0.42	45	35	True
3	04-01-1948	0.31	45	34	True
4	05-01-1948	0.17	45	32	True

```
[4]: data.isna().sum()
```

```
[4]:
```

DATE	0
PRCP	3
TMAX	0

```
TMIN      0  
RAIN      3  
dtype: int64
```

```
[5]: data[data["PRCP"].isna()]
```

```
[5]:          DATE  PRCP  TMAX  TMIN RAIN  
18415  02-06-1998    NaN    72    52  NaN  
18416  03-06-1998    NaN    66    51  NaN  
21067  05-09-2005    NaN    70    52  NaN
```

```
[6]: data[data["RAIN"].isna()]
```

```
[6]:          DATE  PRCP  TMAX  TMIN RAIN  
18415  02-06-1998    NaN    72    52  NaN  
18416  03-06-1998    NaN    66    51  NaN  
21067  05-09-2005    NaN    70    52  NaN
```

```
[7]: data = data.dropna()
```

```
[8]: data.isna().sum()
```

```
[8]: DATE      0  
PRCP      0  
TMAX      0  
TMIN      0  
RAIN      0  
dtype: int64
```

```
[9]: data.RAIN.value_counts()
```

```
[9]: RAIN  
False     14648  
True      10900  
Name: count, dtype: int64
```

```
[10]: from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

```
[11]: data["RAIN"] = le.fit_transform(data["RAIN"])
```

```
C:\Users\prana\AppData\Local\Temp\ipykernel_14568\66160035.py:1:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <https://pandas.pydata.org/pandas->

```
docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy  
data["RAIN"] = le.fit_transform(data["RAIN"])
```

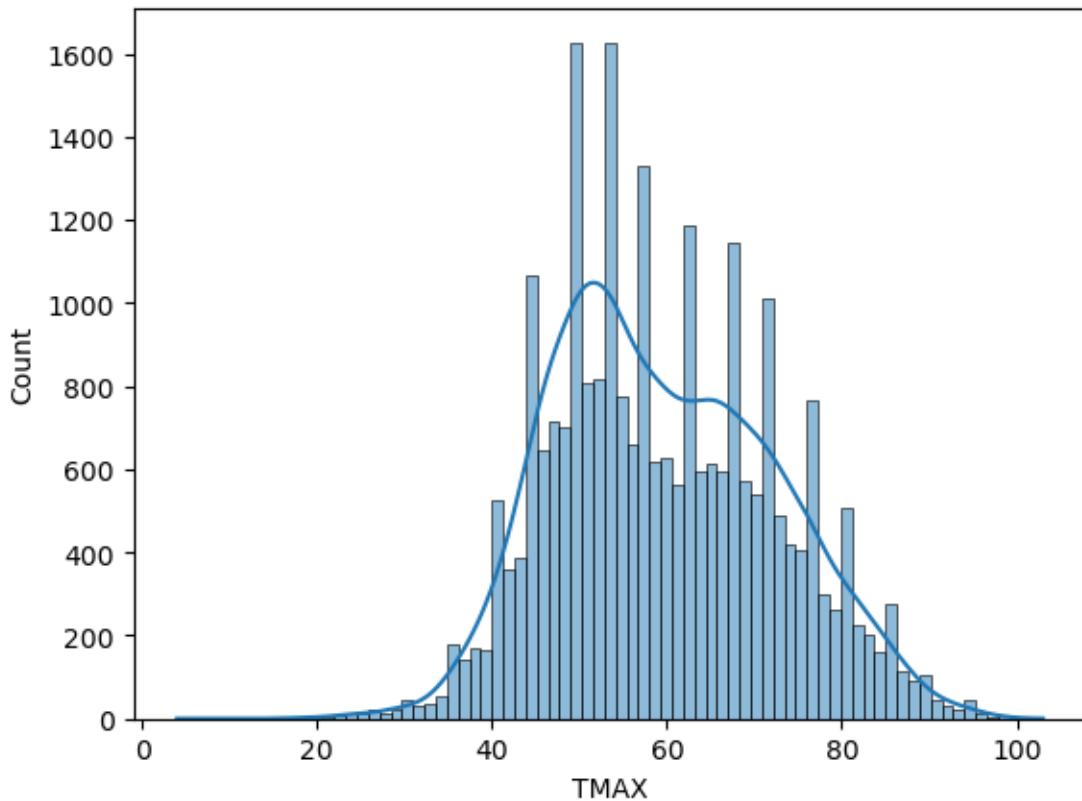
```
[12]: data["RAIN"].dtype
```

```
[12]: dtype('int32')
```

```
[13]: import seaborn as sns
```

```
[14]: #pip install seaborn --upgrade
```

```
[15]: sns.histplot(data["TMAX"], kde = True)  
plt.show()
```



```
[16]: sns.distplot(data["TMAX"], kde = True)  
plt.show()
```

```
C:\Users\prana\AppData\Local\Temp\ipykernel_14568\2605151143.py:1: UserWarning:
```

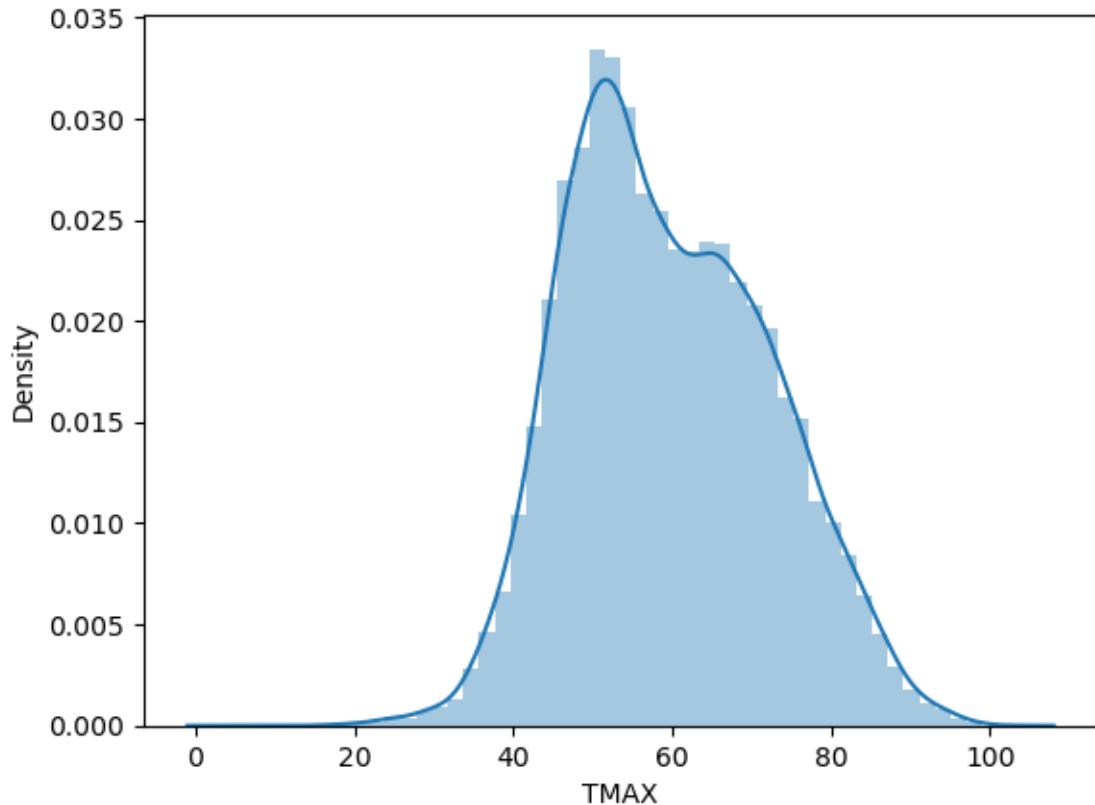
```
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

```
Please adapt your code to use either `displot` (a figure-level function with
```

similar flexibility) or `histplot` (an axes-level function for histograms).

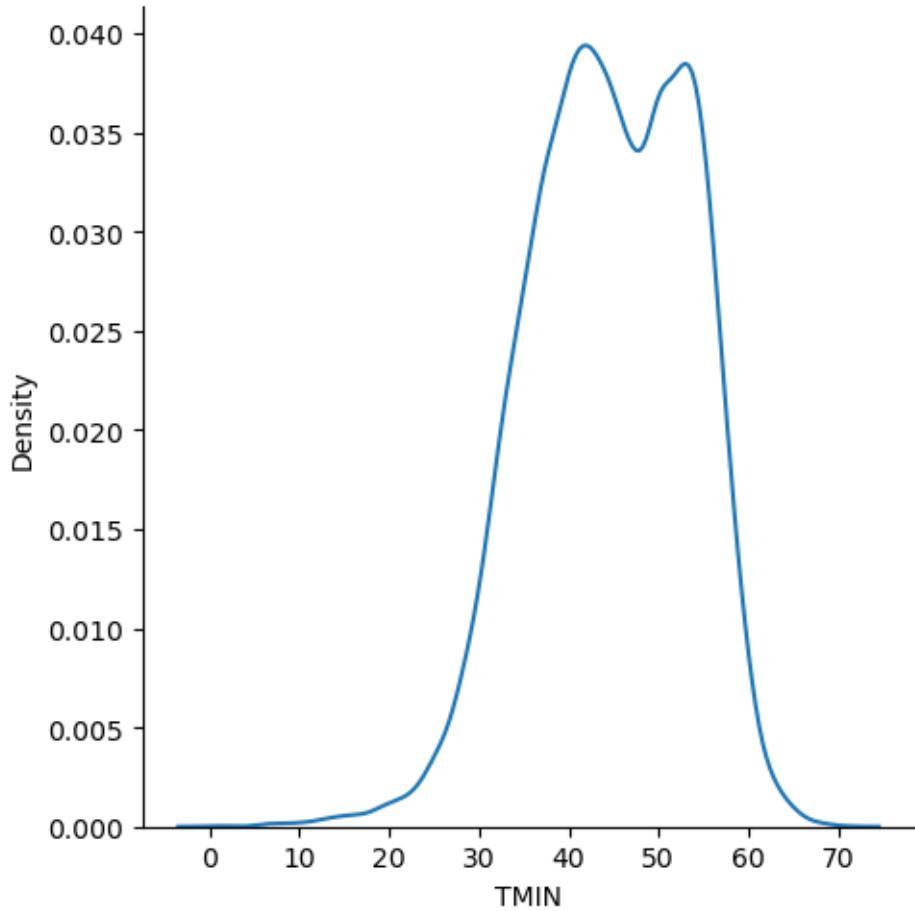
For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(data["TMAX"], kde = True)
```



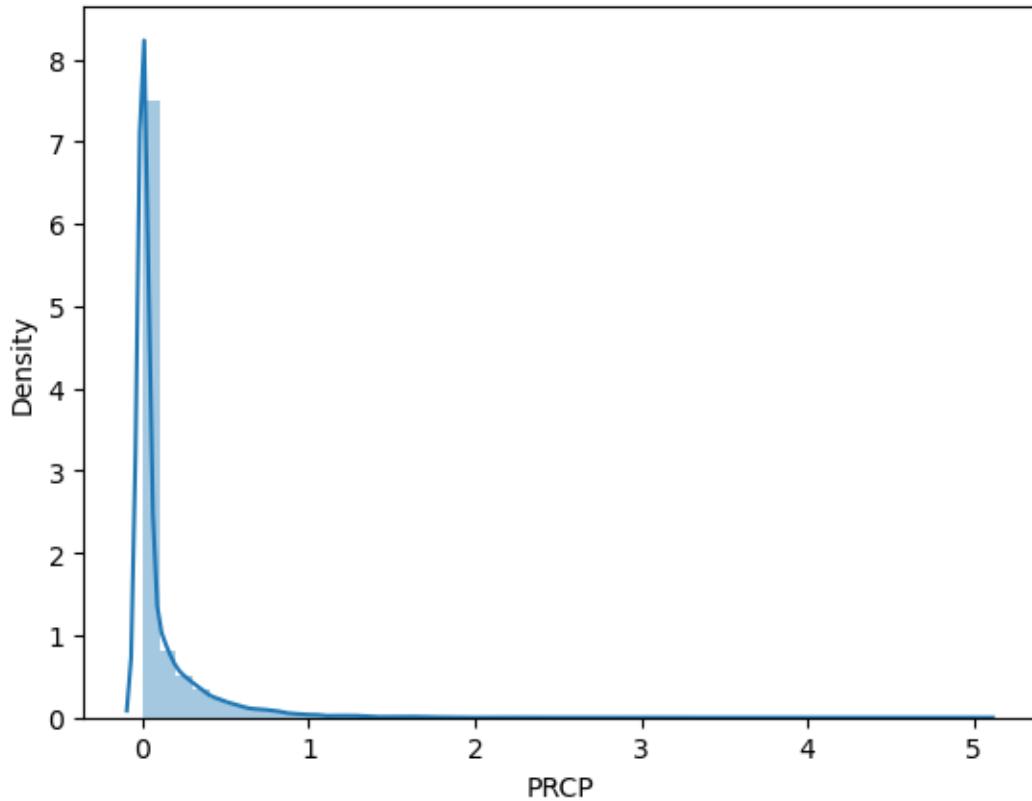
```
[17]: sns.distplot(data["TMIN"], kind ="kde" )
plt.show()
```

```
C:\Users\prana\anaconda3\New folder\Lib\site-packages\seaborn\axisgrid.py:118:
UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)
```



```
[18]: import warnings
warnings.filterwarnings("ignore")

sns.distplot(data["PRCP"])
plt.show()
```

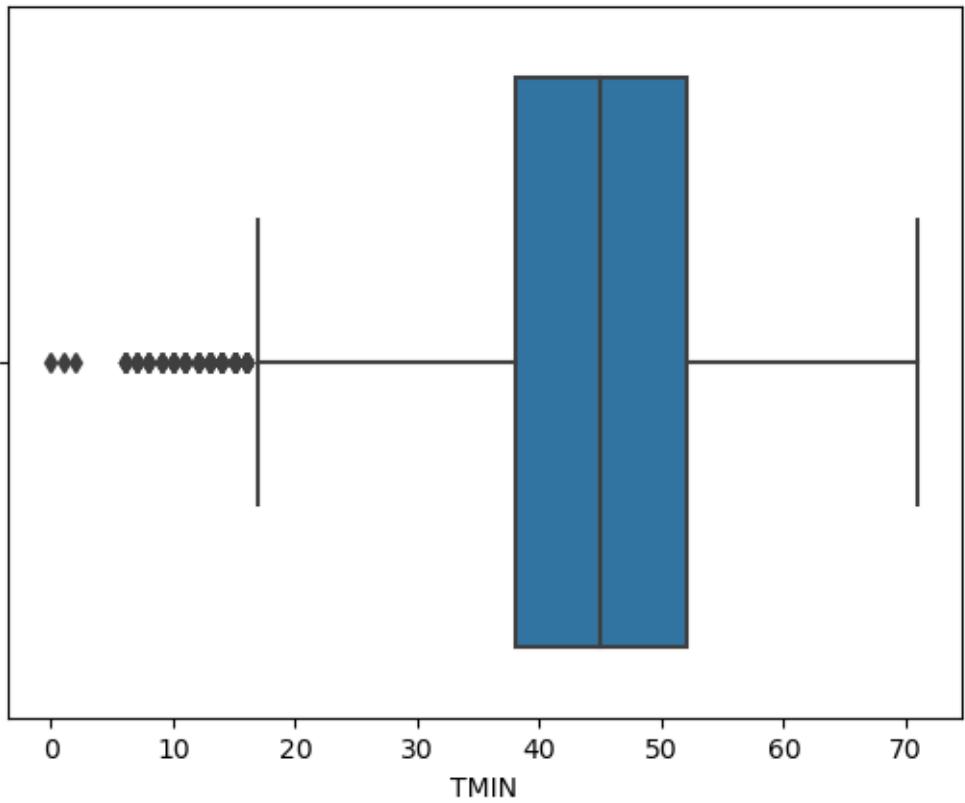


```
[19]: data.describe()
```

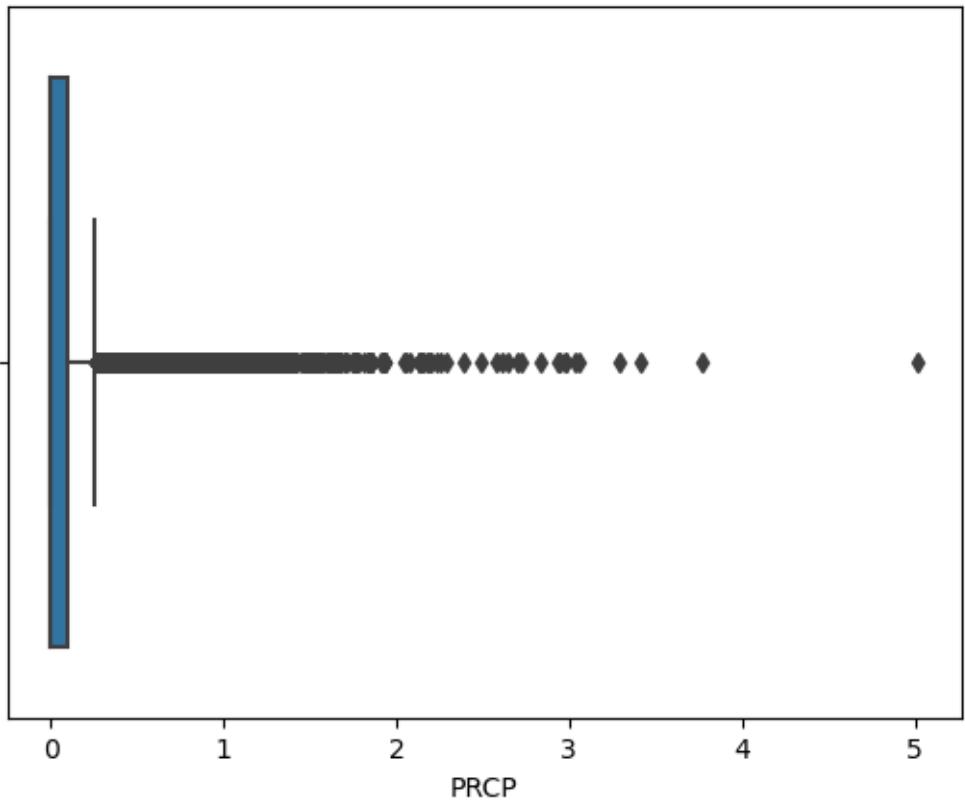
```
[19]:
```

	PRCP	TMAX	TMIN	RAIN
count	25548.000000	25548.000000	25548.000000	25548.000000
mean	0.106222	59.543056	44.513387	0.426648
std	0.239031	12.773265	8.893019	0.494600
min	0.000000	4.000000	0.000000	0.000000
25%	0.000000	50.000000	38.000000	0.000000
50%	0.000000	58.000000	45.000000	0.000000
75%	0.100000	69.000000	52.000000	1.000000
max	5.020000	103.000000	71.000000	1.000000

```
[22]: sns.boxplot(x=data["TMIN"])
plt.show()
```



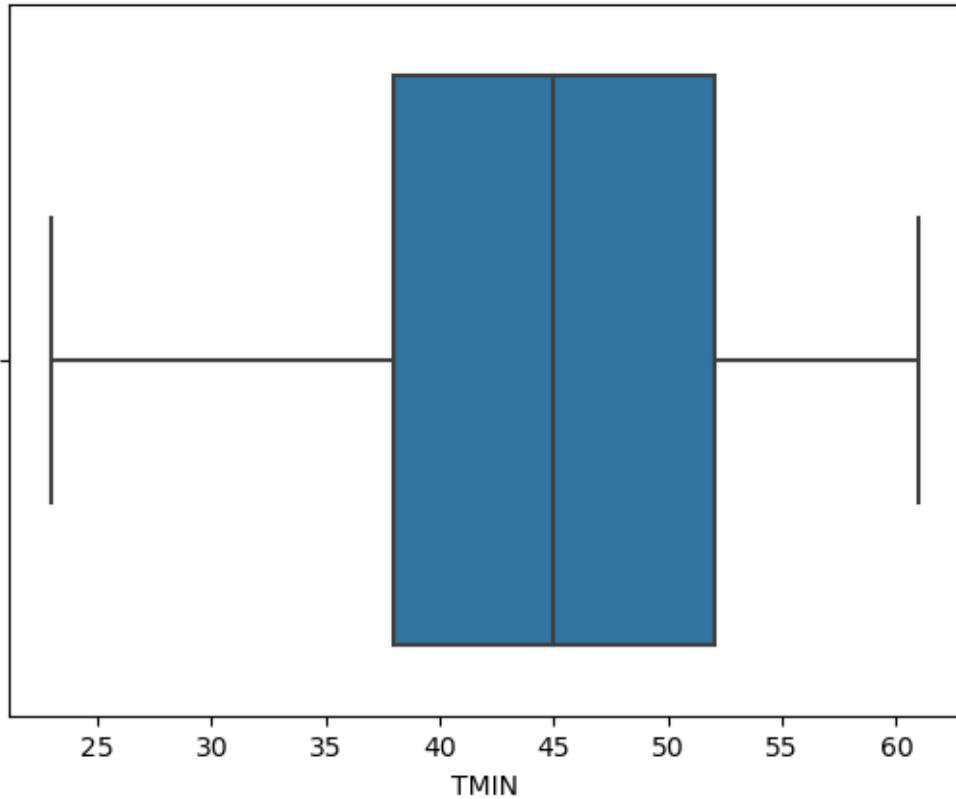
```
[21]: sns.boxplot(x=data["PRCP"])
plt.show()
```



```
[23]: upper_limit = data["TMIN"].quantile(0.99)
lower_limit = data["TMIN"].quantile(0.01)
```

```
[24]: data["TMIN"] = np.where(data["TMIN"] < lower_limit, lower_limit, data["TMIN"])
data["TMIN"] = np.where(data["TMIN"] > upper_limit, upper_limit, data["TMIN"])
```

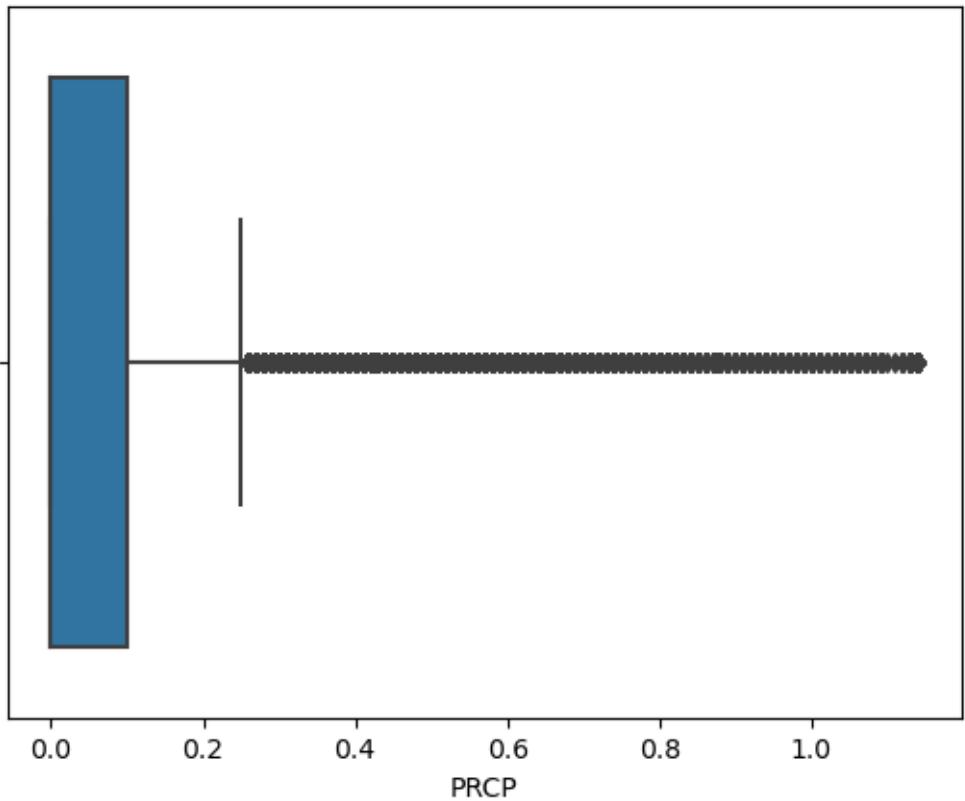
```
[25]: sns.boxplot(x=data["TMIN"])
plt.show()
```



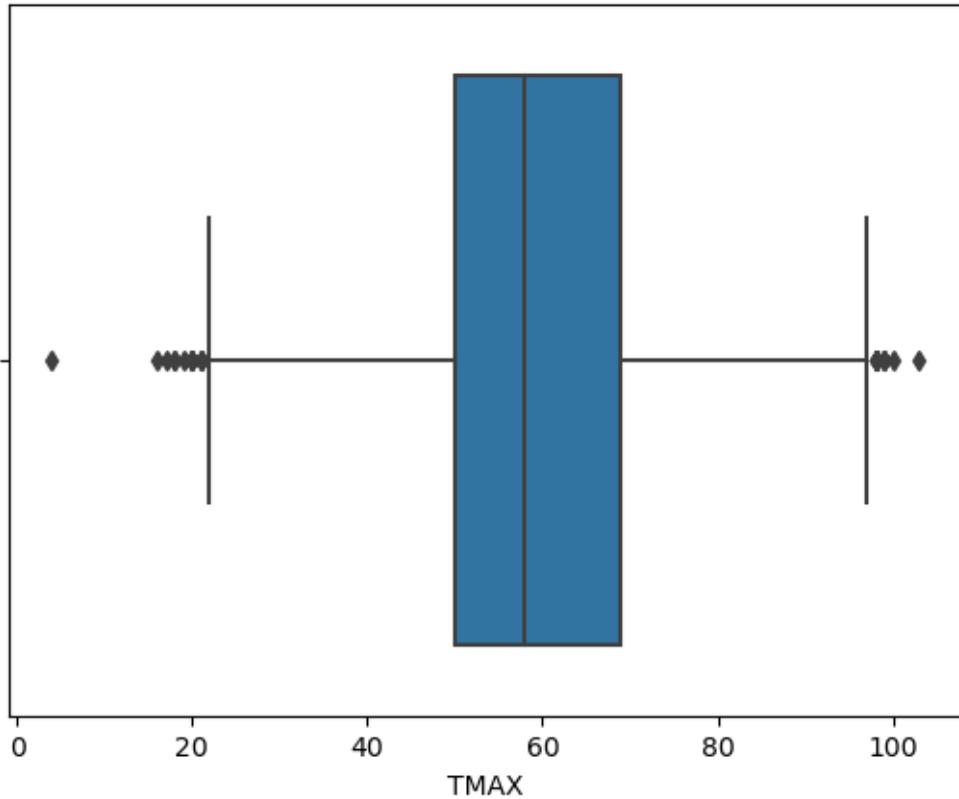
```
[26]: upper_limit = data["PRCP"].quantile(0.99)
lower_limit = data["PRCP"].quantile(0.01)
```

```
[27]: data["PRCP"] = np.where(data["PRCP"] < lower_limit, lower_limit, data["PRCP"])
data["PRCP"] = np.where(data["PRCP"] > upper_limit, upper_limit, data["PRCP"])
```

```
[28]: sns.boxplot(x=data["PRCP"])
plt.show()
```



```
[29]: sns.boxplot(x=data["TMAX"])
plt.show()
```



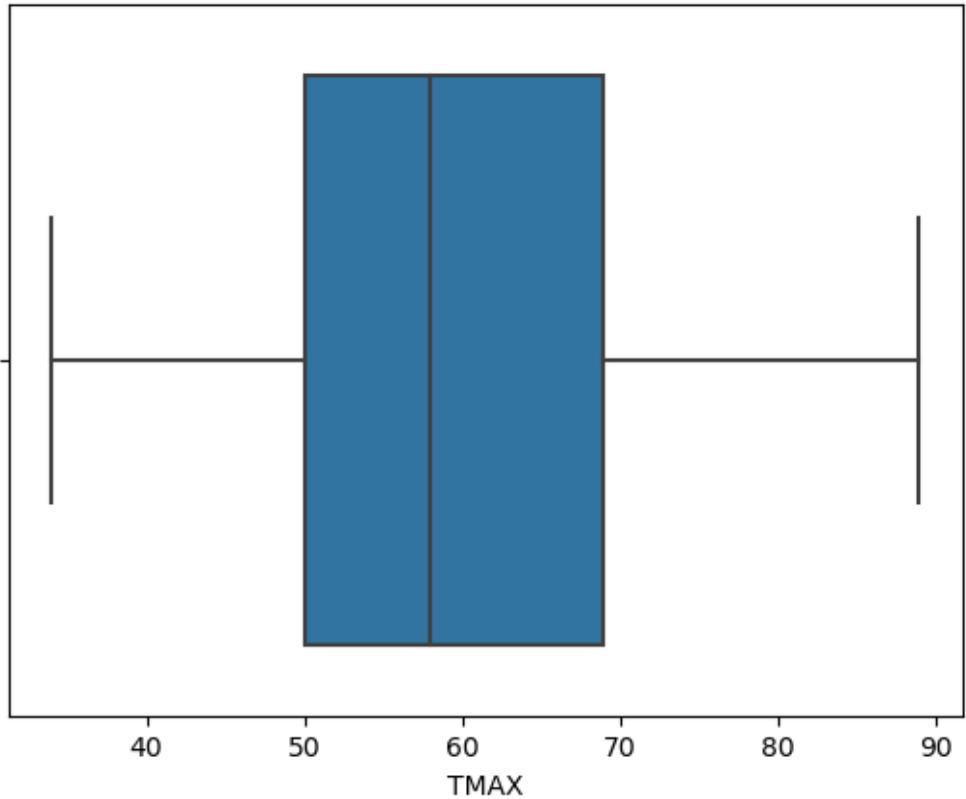
```
[37]: data.shape
```

```
[37]: (25548, 5)
```

```
[30]: upper_limit = data["TMAX"].quantile(0.99)
lower_limit = data["TMAX"].quantile(0.01)
```

```
[31]: data["TMAX"] = np.where(data["TMAX"] < lower_limit, lower_limit, data["TMAX"])
data["TMAX"] = np.where(data["TMAX"] > upper_limit, upper_limit, data["TMAX"])
```

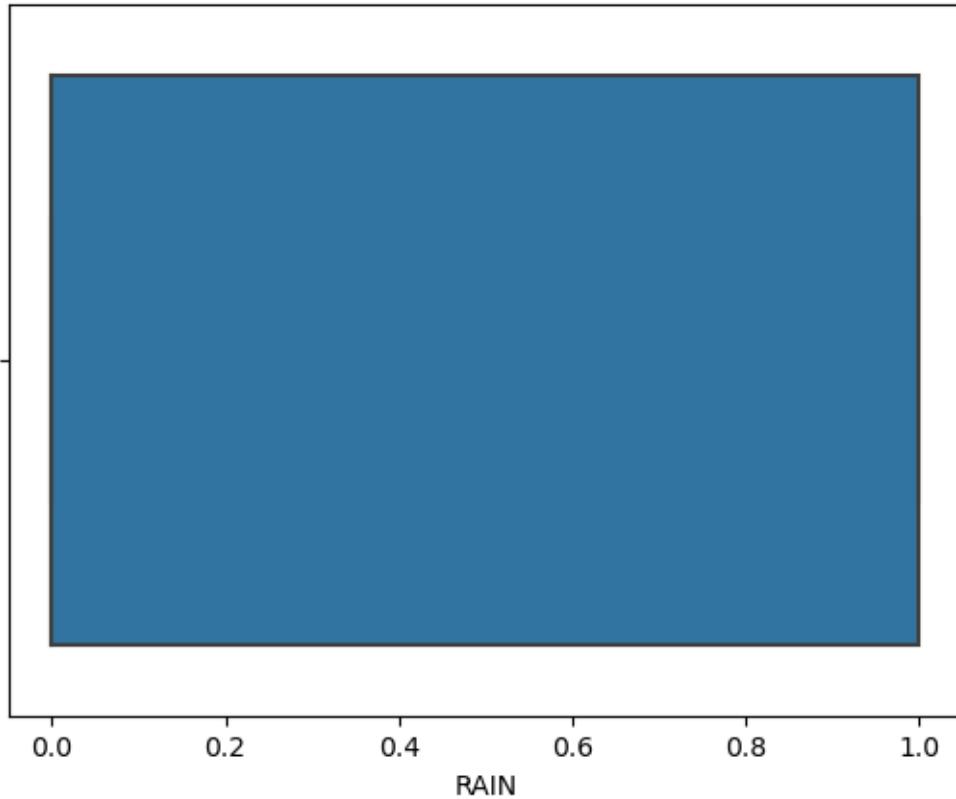
```
[32]: sns.boxplot(x=data["TMAX"])
plt.show()
```



```
[36]: data.shape
```

```
[36]: (25548, 5)
```

```
[33]: sns.boxplot(x=data["RAIN"])
plt.show()
```



```
[35]: data.shape
```

```
[35]: (25548, 5)
```

```
[39]: data = data.drop(data[(data["TMAX"] > 97) | (data["TMAX"] < 21)].index)
```

```
[41]: data = data.drop(data[data["TMIN"] < 17].index)
```

```
[42]: data = data.drop(data[data["PRCP"] > 0.2].index)
```

```
[44]: data.shape
```

```
[44]: (21306, 5)
```

```
[45]: X = data.drop(["RAIN", "DATE"], axis=1)
y = data["RAIN"]
```

```
[46]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.70, random_state=42)
```

```
[47]: #Creation of model
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()

[48]: #Model has learned from the training data
lr.fit(X_train,y_train)

[48]: LogisticRegression()

[49]: pred = lr.predict(X_test)

[50]: X_test.head()

[50]:      PRCP    TMAX   TMIN
13132  0.00  43.0  35.0
25234  0.00  43.0  29.0
11960  0.20  63.0  56.0
23422  0.07  45.0  38.0
16679  0.00  79.0  53.0

[51]: pred

[51]: array([0, 0, 1, ..., 0, 1, 0])

[52]: #Actual value
y_test

[52]: 13132      0
25234      0
11960      1
23422      1
16679      0
..
23717      1
6203       1
16128      0
25353      1
10251      0
Name: RAIN, Length: 6392, dtype: int32

[53]: from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_test,pred)
accuracy

[53]: 0.905819774718398

[54]: lr.predict_proba(X_test)
```

```
[54]: array([[6.72968884e-01, 3.27031116e-01],  
           [8.56642894e-01, 1.43357106e-01],  
           [4.30985713e-04, 9.99569014e-01],  
           ...,  
           [9.78767931e-01, 2.12320687e-02],  
           [1.28297250e-01, 8.71702750e-01],  
           [7.13186420e-01, 2.86813580e-01]])
```

```
[55]: from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test,pred)  
cm
```

```
[55]: array([[4367, 19],  
           [ 583, 1423]], dtype=int64)
```

```
[56]: #TP=4367, TN=1423, FP=19, FN=545
```

```
[58]: #Prediction on new data  
X_new=[[2.2,45,20]]  
pred_new=lr.predict(X_new)  
pred_new
```

```
[58]: array([1])
```

```
[59]: lr.predict_proba(X_new)
```

```
[59]: array([[0., 1.]])
```

```
[ ]:
```